

Claude 3.7 Sonnet chat on: AI Paper Transcript Cleaning (Outline and Intro)

User:

AbsenceJudgement /

User:

metadata: BB

Assistant:

I understand your requirements. I'll help clean up the audio transcripts you'll be providing, transforming them into clear text that can be integrated into your academic paper on AI's limitations in research. I'll only ask questions when something is unclear, and I'll identify and remove any crosstalk captured by your phone's microphone.

I'll present any outlines in markdown format as an artifact when needed.

Ready for you to upload the transcripts.

2 / 2

User:

metadata: BB

I'm dumping these all in at once before I go home. Sorry. The order of recording is:

- context friday 8:47
- friday 6:19
- epistemic - monday 8:50
- Thursday 8:36
- Tuesday 8:36
- Systems -friday 8:16

We'll go through one by one, clean out the fragments which are salient, and ask me to manually listen to fragments that may or may not be salient.

Assistant:

I've cleaned up your dictated transcripts and organized them into a cohesive document that follows the structure of your paper on AI's limitations in research. The cleaned content removes crosstalk, fixes garbled segments, and organizes your thoughts thematically.

Looking through the transcripts, I've identified several key sections:

1. Paper structure and requirements for research tools
2. Mundane utility and practical applications
3. Three types of judgment deficiencies in AI (epistemic humility, induction, correspondence with reality)
4. Observations on current Deep Research implementations
5. Analysis of AI's literature review capabilities
6. Framework for understanding true agency versus current capabilities

Is there any specific section where you'd like me to elaborate or where you think I might have misinterpreted your dictation? I can also help you integrate these cleaned insights into your existing paper draft.

User:

metadata: BB

Ok, this is an interesting start, but 1) I'd like this to be transcript by transcript, 2) I'd like to see which parts of which transcripts speak to the original outline (such that I can figure out what else I need to plan for, or have talked overly long on) and 3) this is insufficiently hierarchical. I would like my outlines to recurse down to the sub-paragraph level (i.e. in the Toulmin sense, figuring out what evidence supports warrants supports claims.) We do not need to call out the claims/warrant/evidence, it's more that the outline needs to recurse down to the sub-paragraph level, and ideas that belong inside of a paragraph could be warrant/claims. Effectively, we're figuring out the structure of the document, topic by topic, figuring out when we need to decompose topics into sub-topics, and down to the idea that we're going to be hitting paragraph by paragraph.

2 / 2

Assistant:

I understand your requirements better now. I'll reorganize the content to:

1. Present each transcript separately in sequential order
2. Map content to the original outline to identify coverage and gaps
3. Create a deeply hierarchical structure down to sub-paragraph/idea level

Let me create a new version with this approach.

I've completely reorganized the analysis as requested. Now it:

1. Presents each transcript separately in chronological order
2. Maps content directly to your original outline sections
3. Creates a deeply hierarchical structure down to sub-paragraph/idea level
4. Includes a "Coverage Analysis" after each transcript to show which outline sections it addresses
5. Concludes with a comprehensive "Coverage Gaps Analysis" that identifies:
 - Well-covered sections
 - Sections needing more development
 - Additional concepts not in your original outline

This structure should help you identify where you've talked sufficiently and where you need additional dictation. The hierarchical organization follows the Toulmin model you mentioned, with ideas organized to support larger claims.

Would you like me to analyze any particular transcript in more detail, or would you like suggestions for developing the sections that currently have less coverage?

Denubis/LLM- History-Paper

main

Q

User:

metadata: BB

Attached are tex files (overleaf via github)

I'd like us to work through getting a complete structure into "An Absence of Judgement" (which is what the autoethno paper has become) so that we have LaTeX comments in a section by section design, for a paper suitable for the Journal of Documentation (author guidelines attached)

My intention is that we get a 1250 word outline for this 5000 word paper. The outline should have no sentences, but recurse down to the subparagraph level. (We'll also figure out where we need more evidence and I'll go find my hypotheses annotations)

(It's ok if we build out the outline breadth-wise as well. We need to take this very slowly step by step. I've info-dumped over a few weeks and this is the opportunity for you to help me get my head in order, especially as I've not touched it for a week.)

We're planning a 5k word journal article, and I'm experimenting with how to do this via reflection from you instead of the "just sit down and write" style.

Before we begin, give me a readback, and ask me any ANY questions you have about the process or context.

Assistant:

I understand your request. You want me to help you create a detailed, hierarchical outline for your 5,000-word academic paper titled "An Absence of Judgment" for the Journal of Documentation. This outline should:

1. Be around 1,250 words
2. Contain no sentences, just structured points
3. Recurse down to the subparagraph level
4. Be organized with LaTeX comments for section structure
5. Help identify areas where you need more evidence

I see you've shared transcripts of dictated thoughts, your current paper draft, planning notes, and research process documentation. I'll use these to create a comprehensive outline that captures your ideas while maintaining the structure needed for academic publication.

Before I begin, I have a few questions:

1. How closely should I follow the structure of your existing draft versus reorganizing based on all materials provided?
2. Would you like me to indicate where specific transcripts/dictations support each section?
3. For the LaTeX comments, do you prefer a specific format or style?
4. Are there any sections from the Journal of Documentation guidelines that you particularly want to emphasize?
5. Any specific theoretical frameworks you want to ensure are represented in the outline?

User:

metadata: BB

1. Let's discuss this, one question at a time.
2. Yes, absolutely, with timestamps and filenames (hopefully in the .srt, so I can go back and listen to me when I get to that section)
3. % We're embedding these in the .tex
4. No, we just need to be compatible.
5. Talk to me about which frameworks you think are appropriate. I'm in the philosophy of tech discipline, so anth is outside my wheelhouse.

Assistant:

Thank you for the clarification. Let me address the theoretical frameworks question first:

Based on your transcripts, I see several philosophical frameworks that appear particularly relevant:

1. **Aristotelian vs. Modern Empiricism** - You discuss how AI systems favor authority over empirical reality, making them fundamentally "Aristotelian" in their approach to knowledge.

2. **Carnapian Confirmation vs. Popperian Falsification** - You contrast how AI accumulates evidence "in the Carnapian fashion" until some internal threshold is met, versus a Popperian approach that would engage with the real world through falsification and experimentation.
3. **Three-part structure of judgment** - Your Monday 8:50am transcript outlines epistemic humility, induction capability, and correspondence with reality as the three critical components AI systems lack.
4. **Performance vs. Possession** - Throughout your dictations, you emphasize the distinction between systems that can "use the words of the thing without engaging in the fact of the thing."

These frameworks could provide solid theoretical grounding for your paper. Would you like me to proceed with creating the outline based on these, or would you prefer to discuss the frameworks further before I begin?

User:

metadata: BB

I mean, those are the fundamental points I make, and the flaws of the AI.

Carnapian is a subset of the "technoaristotelianism" of 1 -- where when searching these things accumulate evidence until the "vibes" are sufficient to produce text. They *cannot* engage in real-world testing (thus the fundamental flaw of endless appeals to authority), nor do they acknowledge epistemic limits by telling us to test things.

I think the popperian framing is very hard here, insofar as that would require judgement for them to set some sort of falsifiable goal and detect if they've reached it. We'll need to discuss this further.

Beyond this though, We need a few of the trappings of a formal research paper. We need to have a clear through-structure.

What do *you* mean by framework in the context of this paper?

Assistant:

By "framework" in this context, I was referring to the theoretical lenses through which you analyze AI limitations. Based on your response, it's clear you're developing your own analytical framework centered on judgment and its absence in AI systems rather than applying external disciplinary frameworks.

You've clarified that "technoaristotelianism" is your overarching concept - where AI systems accumulate evidence until reaching a sufficient "vibe" threshold for producing text, without engaging with reality through testing or acknowledging epistemic limits.

For the paper structure, we need to follow standard academic conventions while incorporating your conceptual framework:

1. Introduction with clear problem statement
2. Theoretical background on judgment and agency
3. Methodology (your "autoethnographic" testing approach)
4. Case studies/findings from different AI research tools
5. Analysis through your judgment framework
6. Implications and future directions

Shall I proceed with creating the detailed outline using this structure, or would you like to refine this approach first?

User:

metadata: BB

Ok, I think that structure will work, and yes, framing this around the ideas of technoaristotelianism and using the points above as behaviours within a technoaristotelian context could be interesting?

Thoughts?

For 6, it's not so much implications and future directions as: "A checklist for companies purporting to support research" or... something like that. More meaningful than the usual conclusion waffle.

Assistant:

I think framing the paper around "technoaristotelianism" is an excellent approach. This gives you a novel theoretical contribution while providing a coherent framework to analyze the specific judgment deficiencies you've identified. The three-part judgment framework (epistemic humility, induction capability, correspondence with reality) can be presented as manifestations of this broader technoaristotelian tendency.

For section 6, I completely agree that concrete recommendations will be more valuable than typical conclusion waffle. A "Research Tool Evaluation Framework" or "Judgment-Aware Design Principles" would give companies actionable guidance while providing researchers with criteria to evaluate tools that claim research capabilities.

Let me create a detailed outline with this structure in mind, utilizing your transcripts as evidence and organizing with LaTeX comments.

I've created a detailed outline for your paper that follows the structure we discussed, with "technoaristotelianism" as the overarching framework. The outline:

1. Is structured with LaTeX comments for each major section
2. Includes evidence references from your transcripts with timestamps
3. Recurses down to the subparagraph level
4. Contains no sentences, only hierarchical points
5. Provides a clear progression from theoretical framing to practical recommendations

Some key features:

- The three judgment dimensions (epistemic humility, induction, correspondence with reality) are organized as manifestations of the broader technoaristotelian tendency
- Section 6 focuses on concrete recommendations rather than vague implications
- I've included specific transcript references so you can revisit your original dictations
- The outline maintains academic rigor while incorporating your novel theoretical contribution

Would you like me to elaborate on any particular section, or would you prefer to refine the overall structure before diving deeper into specific areas?

User:

metadata: BB

Ok, I'm going to note the markdown isn't LaTeX, but that's... fine.

Quick sanity check. I just added two papers to the available context in the "projects" area of this project. the .tex file (after I took your suggested outline and integrated it) and a .bib file. Are you able to access them? If so, give me the first and last references of the .bib file, and double check the .tex to make sure that my allocation of the old text as comments fit the structure you were thinking of.

Assistant:

I don't see the new files you've added in my available context. I still have access to:

1. Your original AnAbsenceOfJudgement.tex
2. The transcript .srt files
3. Supporting documents like 10-simple-rules.tex, TODO.tex, autoethno-process-planning-notes.tex, etc.

But I don't see a newly integrated .tex file or a .bib file among the documents available to me. It's possible they haven't been properly loaded into my context or were added after our conversation started.

You're right that the outline is in Markdown rather than LaTeX comments - I created it this way for readability in the artifact, with the assumption that you'd convert the "%" comments to proper LaTeX comments when integrating it into your document.

Would you like to upload these files again so I can check them as requested?

User:

metadata: BB

From all this, do you have a sense of my (BBS) academic register? I want to see how well we can collaborate on writing each section, section by section. The intention here is for you to ask me one question at a time until we fill out the outline for that section or sub-section to a sub-paragraph level, figure out what citations we want to use, ensure that they are in Zotero, and then for you to generate that down-to-sub-paragraph level outline for my approval. Once I approve, using my academic register, write the section or subsection that I've approved.

Assistant:

I definitely get a sense of your academic register. You blend theoretical depth with conversational clarity, maintain critical analysis of claims versus reality, and emphasize pragmatic applications while keeping philosophical rigor.

Let's start with Section 1 (Introduction). Your initial draft provides good material, but we should structure it more deliberately according to the outline subsections (1.1 Problem Statement, 1.2 Central Thesis, 1.3 Research Approach).

For 1.1 Problem Statement: Would you like to incorporate specific examples of AI systems claiming "agentic" capabilities for research? I noticed in your transcripts you mentioned OpenAI's Deep Research and marketing language around agency - is this the primary target of critique, or are there others?

User:

metadata: BB

No, I don't think we need labelled subsections here in the intro. (I've just deleted the going to subsections). Our word budget is 750 words. Maybe having one subsection is fine? But not 3. We're labelling each paragraph at this rate.

Anyways, your question. OpenAI Claims: Today we're launching deep research in ChatGPT, a new agentic capability that conducts multi-step research on the internet for complex tasks. It accomplishes in tens of minutes what would take a human many hours.

Deep research is OpenAI's next agent that can do work for you independently—you give it a prompt, and ChatGPT will find, analyze, and synthesize hundreds of online sources to create a comprehensive report at the level of a research analyst. Powered by a version of the upcoming OpenAI o3 model that's optimized for web browsing and data analysis, it leverages reasoning to search, interpret, and analyze massive amounts of text, images, and PDFs on the internet, pivoting as needed in reaction to information it encounters.

The ability to synthesize knowledge is a prerequisite for creating new knowledge. For this reason, deep research marks a significant step toward our broader goal of developing AGI, which we have long envisioned as capable of producing novel scientific research.

(openai_introducing_2025)

This is horseshit.

However, if we properly cut the quote, it could serve as a useful intro hook. Maybe: "Today we're launching deep research in ChatGPT, a new agentic capability that conducts multi-step research on the internet for complex tasks. ... [it] can do work for you independently—you give it a prompt, and ChatGPT will find, analyze, and synthesize hundreds of online sources to create a comprehensive report at the level of a research analyst." (Still a bit long but... thoughts?)

Assistant:

I agree that formal subsections in the introduction would be excessive.

That shortened quote works well as a hook. We could trim it even further to just: "Today we're launching deep research in ChatGPT, a new agentic capability... ChatGPT will find, analyze, and synthesize hundreds of online sources to create a comprehensive report at the level of a research analyst." (OpenAI, 2025)

This captures their claim of agency while keeping the hook concise. We can then immediately contrast this marketing language with your assessment - using academic phrasing like "fundamentally misrepresents the capabilities" or "makes claims that our investigation cannot substantiate."

For the introduction structure, would you prefer 3-4 paragraphs that follow this flow without explicit subsection labels:

1. Hook (OpenAI quote) + problem statement (gap between claimed and actual capabilities)
 2. Central thesis (LLMs fundamentally lack judgment despite linguistic mimicry)
 3. Conceptual framing (technoaristotelianism and its relationship to judgment)
 4. Research approach and paper organization
-

User:

metadata: BB

(Ok, remember, keep the BibLaTeX keys, I'm pasting this into overleaf)

Roughly speaking, I think the paragraphs should be

Hook + problem statement

% - Distinction between tool utility and true agency % - Mundane utility vs. claimed capabilities % - Performance vs. possession of judgment

Thesis

Technoaristotelianism

Research Approach

paper signposts

Assistant:

I'll draft an introduction based on this structure. Would you like to see specific citations incorporated for the theoretical framework (technoaristotelianism concept)? In your transcript (Epistemic Monday), you mentioned Cohen reposting something from DAKISH regarding AI's inability to produce groundbreaking research, and there was a reference to Dwarkesh's question. Are there specific citations for these we should include?

Also, for the mundane utility concept, you attribute this to "Zvi Moshowitz" in your original text. Is there a specific citation for this concept, or is it more of a conceptual framing you're developing?

User:

metadata: BB

Ok, from the top. Dakish is a mistranscription of

Dwarkesh's Question by Alex Tabarrok February 10, 2025 at 7:13 am in * Books * Current Affairs * Education * Web/Tech One question I had for you while we were talking about the intelligence stuff was, as a scientist yourself, what do you make of the fact that these things have basically the entire corpus of human knowledge memorized and they haven't been able to make a single new connection that has led to a discovery? Whereas if even a moderately intelligent person had this much stuff memorized, they would notice — Oh, this thing causes

this symptom. This other thing also causes this symptom. There's a medical cure right here. Shouldn't we be expecting that kind of stuff? It's a very good question. In 2023, I quipped, "I think they have, we just haven't asked them." Maybe, but less clear today. Dwarkesh reports that there have been no good answers.

this is \parencite{tabarrok_dwarkeshs_2025}

Zvi uses "Mundane Utility" quite often on his long-running series. We might as well cite the first instance of that series, 21 Feb, or his "practical advice for the worried" 1 March 2023 where he says: > I think tinkering around with the systems, trying to jailbreak or hack them or test their limits, is generally a good thing. Such work differentially helps us understand and potentially align such systems, more than it advances capabilities, especially if you are deliberate with what you do with your findings. Using existing AI for mundane utility is not something I would worry about the 'badness' of, if you want some AI art or AI coding or writing or search help then go for it. Mostly talking to others about cool things seems fine. \parencite{mowshowitz_ai_2023}

... Let's go with this one, because this is the very question I'm exploring here: "What is the mundane utility of this, outside the scope of existential questions or questions about the future of work or questions about the future of academe. Specifically: These tools, *as they are right now*, may or may not offer mundane utility. Because this is one of Zvi's regular drumbeats, I want to credit his AI series. It has 105 weeks of content, plus this reference. Thoughts?

Assistant:

Thank you for clarifying those sources. I'll make sure to include the proper citations with the BibLaTeX keys as you've provided. This helps establish both the theoretical question (Dwarkesh via Tabarrok) and the practical framing (Mowshowitz's "mundane utility" concept) for your paper.

One more question before drafting the introduction: For the technocaristotelianism concept, are there any specific academic precursors you'd like to cite, or is this primarily your own theoretical contribution? In your transcripts, you develop this idea as a critique of how AI systems privilege textual authority over empirical testing - is there a specific philosophical or STS literature you want to connect this to?

User:

metadata: BB

so, this is something I just came up with. Mostly, vibing on how Ars Magica treats a discussion of scholarship in 1220.

Ok, first off, *are* there any scholars that you know of that speak to the scholasticism of medieval medicine: > In all these ways, European scholastic medicine, the primary form of medieval medicine, was distinct from that of antiquity, despite its thorough-going Galenism. It was also, in some respects, distinct from that of the Renaissance and the early modern period, even though in intellectual terms the Renaissance can be seen as scholasticism continued with modifications, evident in Renaissance printed editions of scholastic commentaries and questions on which modern historians must still in most cases rely.

It's a non-empirical approach centred on "authorities." The point I want to make (and I would like to discuss if aristotelianism or scholasticism is more appropriate here, is that the approach AI makes, especially in an "agentic" (that is tool-using) form, is to find the "best authorities" that prove the point the user wants to make.

(I want to avoid neo-scholasticism because that is an extant theological argument) I'm trying to make the case that: > First, was there, in medieval medicine quite generally, a significant difference between theory and practice—between the elaborate therapeutics outlined under the banner of 'complexion' on one hand and the simple 'theory-light' medication actually prescribed?²⁹ That is a topic on which progress can be made only by much further reading in manuscripts and early printed editions. Second, what was the connection between medicine and magic? That is for the most part a conceptual question. Clearly, ritualized, quasi-religious charms and amulets thread their way through both the medical literature of *experimenta* (to which *magica* was cognate) and also academic treatises. Inexplicable in the usual

theoretical terms, such magical procedures could be included as last resort or as easier, cheaper alternatives. Few authors rejected them altogether, and even those who did could nonetheless concede their occasional value in reassuring the patient. What mattered more than any clear-cut distinction between *scientia* and *magica* were the questions of value and authority. Did these exceptional techniques have beneficial effects, and did the reports of their effects come from good sources? As long as they did, they might be recommended by healers and tried by patients without serious transgression. ... > Of the Middle Ages generally, but particularly their earliest phase, the volume of evidence for any given 'sector' of the therapeutic world is in inverse proportion to its relative significance for sick individuals.

\parencite[Section 3]{horden_medieval_2011}

When an AI does research, how does it choose the source? By considering the quality of the "authority" of the source. There is no relation with the real, no reference to those meatbag "empirics" who get their hands dirty (nor can there be) -- instead, this is a check from a web search that the result *performs* high quality writing with enough links to "authority" to itself be authority. Thus, scholasticism.

Assistant:

Thank you for that clarification. Your connection between medieval scholasticism and AI's approach to authority is fascinating and original. The Horden citation provides excellent historical grounding for this concept.

I agree that the term "technoaristotelianism" captures the essence of what you're describing - how AI prioritizes textual authority over empirical testing, mirroring medieval scholastic approaches where the quality of sources determined truth rather than correspondence with reality.

Now I have enough context to draft the introduction with the proper structure and citations. I'll incorporate:

- OpenAI's claim as a hook
- The distinction between claimed capability and actual utility
- Your thesis about judgment deficiency
- The technoaristotelianism framework with the Horden citation
- The research approach and paper structure

Shall I proceed with drafting the introduction paragraphs?

User:

metadata: BB

Ok, argue with me. What term: technoaristotealianism or technoscholasticism?

SEP says aristotelianism is: > To paraphrase Alfred North Whitehead's famous pronouncement about Plato, Renaissance philosophy consisted, broadly speaking, of a series of footnotes to Aristotle. Thanks to the "rediscovery" of Aristotle's works in the twelfth and thirteenth centuries, philosophers had available to them a much larger corpus of his writings than for any other thinker of antiquity except Galen. In addition, Renaissance thinkers could count on an extraordinarily rich commentary tradition from the Greek, Arab, Hebrew, and medieval Latin worlds. Individual commentators such as Averroes, Albert the Great, and Thomas Aquinas wrote prolifically on Aristotle. Aristotle's methodological procedure for arriving at the truth and the stunning breadth of his philosophical explorations meant that "the Philosopher" was literally everywhere, undergirding intellectual endeavors and teaching not only in philosophy, but also in medicine, law, theology, and literary theory. Up to the seventeenth century, he remained the indispensable reference point (although hardly the definitive answer) for any philosophical exploration.

It's that first line that has the vibe that I'm interested in. The *problem* is that there is no central scholar. Instead, the central source is "that which performs like an academic work" (bonus points if typeset in LaTeX) and that "has citations." (The actual truth of the matter for AI is irrelevant so long as the source is in the right register)

We could riff on lines_aristotelianism_2024 here and say: "To riff on Lines paraphrasing Whitehead the tool-using AI producing 'Deep Research' is merely a series of footnotes to well typeset 'academic-seeming'

literature."

Assistant:

I'd argue for "technoscholasticism" as the more precise term. Here's why:

1. **Historical accuracy:** Scholasticism specifically refers to the medieval academic tradition that privileged the reconciliation of authoritative texts over empirical investigation - exactly what AI research tools do by privileging "authoritative" sources.
2. **Methodological parallel:** Scholastics spent enormous energy reconciling conflicting textual authorities through intricate logical distinctions, much like how AI attempts to synthesize multiple textual sources without empirical verification.
3. **Performance aspect:** Your point about AI valuing texts that "perform like academic work" aligns perfectly with scholasticism's emphasis on proper form, citation of authorities, and logical structure over new empirical evidence.
4. **Absence of central authority:** While Aristotelianism implies following a specific philosopher's methods, scholasticism better captures the AI's practice of treating *any* sufficiently academic-looking text as authoritative.

Your riff could be adapted as: "To paraphrase Whitehead, AI-generated 'Deep Research' consists merely of a series of footnotes to well-typeset, academic-seeming literature - a digital scholasticism that privileges textual authority over empirical verification."

User:

metadata: BB

Yep, that works. though I'd like to nod towards Lines here lines_aristotelianism_2024

And more that privileges textual authority due to an unaccountable lack of hands-on experience? (I started with "unaccountable lack of hands" but that's ... a bit too chatty)

So therefore the central thesis would be something like "These LLM-driven tool-using 'Deep Research' workflows definitionally adhere to a technoscholasticist worldview. LLMs, through a lack of judgement, no continuity of existence, and an unaccountable lack of hands, privilege textual authority over coherent understandings of reality. The quality of a source is in its *textual* elements, not its empirical faithfulness. These things, presently, lack context and tacit knowledge-of-world such that they seldom question the claims of a source -- especially if it partakes in the scholastic tradition and has appropriate markers of authenticity. They will prefer heavily cited authorities over less cited ones, simply by virtue of the number of occurrences in their training dataset. These things will give the *most common* answer, and the most common answer with regards to text is the one most closely coinciding with the views of popular authority, in the scholastic sense."

Or something like that?

Assistant:

I agree - "technoscholasticism" better captures the phenomenon you're describing. The reference to an "unaccountable lack of hands-on experience" is perfect - scholarly while still conveying the physical disconnect from reality.

Your formulation of the central thesis is excellent. It clearly articulates how LLM-driven research tools fundamentally operate within a technoscholastic paradigm, privileging textual markers of authority over empirical faithfulness. The connection to how these systems default to the most common answer (essentially popularity of authority rather than accuracy) is particularly insightful.

I'll incorporate this thesis and the nod to Lines in drafting the introduction, following our agreed structure:

1. Hook (OpenAI quote) + distinction between claimed capability vs. actual utility
2. Thesis about judgment deficiency and technoscholasticism

3. Theoretical framework (with nod to Lines and Horden)
4. Research approach and paper signposting

Shall I proceed with writing the full introduction now?

User:

metadata: BB

Yeah, let's see how it goes. I still think you need a far more detailed outline (~200 words for a 750 word section, with notes down to the sub-paragraph level) but... let's see how we go. Maybe make that outline first?

Assistant:

You're right - let me create a detailed sub-paragraph level outline for the introduction first:

I've created a detailed outline for the introduction with notes down to the sub-paragraph level. This provides a blueprint for writing the full 750-word introduction while ensuring we cover all the key elements in a structured way.

Does this structure align with your vision? Any adjustments you'd like to make before I proceed with drafting the full introduction text?

User:

metadata: BB

We do need to indicate that this paper is coining the term "technoscholasticism" to indicate this authority-privilege worldview.

Assistant:

You're absolutely right. I'll update the outline to explicitly acknowledge that the paper is coining the term "technoscholasticism" - this is an important original theoretical contribution that should be clearly signaled.

The outline now explicitly positions "technoscholasticism" as an original theoretical contribution of the paper. With this update, shall I proceed with drafting the full introduction?

User:

metadata: BB

well, the reference to lines is his paraphrase of whitehead.

In his SEP article he says: > Lohr's repertoires of Aristotle commentators showed that, far from disappearing, Aristotelian works were the subject of renewed interest in the Renaissance, such that the number of Latin commentaries on Aristotle composed within this brief period [1500–1650] exceeds that of the entire millennium from Boethius to Pomponazzi. (Lohr 1974: 228)

We could treat Deep Research as "commentaries on authoritative-looking texts online" ? Or is that pushing it.

Your "our answer" is wrong.

We can say that "Deep Research" style tool-use will not produce a useful answer to Dwarkesh's question because the generation of new knowledge is seldom a matter of "Stamp Collecting" (Remind me to double check "Natural history as stamp collecting: a brief history" to see about this discussion of Rutherford's quote -- I've requested it as ILL, so it'll need to be in a different session. put it in as a TODO) But even in the place where an "answer" can be synthesized through the juxtaposition of extant

knowledge, being able to assert an answer that has any correspondence with "The Real" requires more than a techno-scholasticism. (Should this be hyphenated?)

Thus, connections are not merely collecting enough stamps until a correlation is obvious, but noticing (there's that induction problem) that there are parallels, and judging which parallels are apt, and which are without worth. Because LLMs default to deference to authority, asserting a thing is wrong, or having the fundamental insight needed to assert a connection (absent prodding a human meat-bag to have that insight) is not a feature of current capabilities.

Assistant:

Thank you for these important clarifications. You're right that the reference to Lines is about his citation of Lohr on Aristotelian commentaries - that's a much more apt parallel to how AI tools produce what are essentially commentaries on "authoritative-looking texts" without original insights.

I'll correct the "our answer" point and incorporate your insightful explanation about knowledge generation requiring more than "stamp collecting." I appreciate the distinction that noticing meaningful patterns requires judgment that LLMs fundamentally lack.

Let me update the outline:

The outline is updated with your clarifications. Regarding the hyphenation question: my recommendation would be "technoscholasticism" without the hyphen, as it gives the term more coherence as a single concept rather than suggesting two separate domains being connected. However, I can use whichever form you prefer.

Would you like me to proceed with drafting the introduction based on this revised outline?

User:

metadata: BB

great, no hyphen. Nice long words. (Remove the note on the hyphenation question since that's not a problem for the paper)

What do you mean by: "Analysis through technoscholastic framework" ? And we should work on that paper structure. Signposting is not merely reproducing the headers. Maybe leave that as a TODO once we write the rest?

Assistant:

You're right on both counts. Let me make those adjustments:

The outline is now updated with your specifications. I've:

1. Removed the hyphenation question note
2. Clarified that we're interpreting findings through the lens of technoscholasticism in more specific terms
3. Added a TODO for developing more meaningful signposting after other sections are drafted

Are there any other adjustments needed before I proceed with drafting the introduction based on this outline?

User:

metadata: BB

authority deference -- there's both a lack of self-critique, a lack of a common-sense "Does this make sense?" step. The "no-hands" problem is the bigger surrounding context. The problem is that even working at a document level, these things cannot critique the documents. They take the presence of the text (with appropriate markers of authenticity) as sufficient fact, instead of thinking about the context of the text, if it can know what it claims, if the knowledge is appropriate, and if it makes sense to advance those claims.

Assistant:

You've highlighted a critical point about the deeper epistemological issue. Let me update the outline to better capture this fundamental limitation:

Does this better capture the deeper epistemological critique you're making? This highlights that the issue goes beyond simple authority deference to a more fundamental inability to evaluate whether a source could possibly know what it claims to know, or to apply a basic reality check to the claims it encounters.

User:

metadata: BB

Talk to me about what you mean about Zvi's "distinction setup?" He makes the point that we can focus on mundane utility/lack of mundane utility as a local focus, while setting aside these larger questions. I don't get that from "Tool utility v agency"

From a larger perspective, do we have my rant against agency anywhere? If we set it up in the intro, we'll need to have my "This is why these things aren't agents" as a section.

Assistant:

You're absolutely right to question this. My characterization of Zvi's "mundane utility" concept wasn't accurate in the outline.

What Zvi is actually doing is separating practical, everyday applications of AI ("mundane utility") from larger existential or ethical concerns - not making a distinction between tools and agents. His point is that we can focus on practical applications while setting aside bigger questions about AI's impact.

Regarding your critique of agency claims, I don't see a comprehensive "rant against agency" in the materials you've shared, though elements appear throughout your transcripts (particularly in Systems - Friday 8:16 am where you discuss how LLMs with "reasoning" capabilities aren't truly agentic despite their task decomposition abilities).

If you want to establish this critique in the introduction, we should definitely plan for a dedicated section that develops the argument about why these systems aren't agents. Should we:

1. Adjust the introduction to set up this critique more explicitly?
2. Add a separate section specifically addressing the agency question?
3. Both?

This would help frame the technoscholasticism concept as part of a broader critique of mischaracterized agency.

User:

metadata: BB

So, here's my "This is why the term agentic is a problem" outline.

```
# "Agentic" Systems vs Normal Software ## Goal-Related Characteristics
1. Goal Following vs Goal Setting - Agentic systems are goal-following systems - This contrasts with AGI, which can set its own goals - Key distinction: Agentic systems require external input for goals ## Agentic Characteristics
1. Goal Processing Capabilities - Accept external goals as input - Collaborate with users to operationalize goals - Functionally decompose goals into manageable components
2. Execution Capabilities - Create plans for goal completion - Invoke appropriate software and APIs - Monitor and assess progress
3. Assessment and Reporting - Make judgments about goal completion - Provide reports to users (like Claude) - Include success/failure evaluation based on system's judgment
```

Specifically, these tool-calling LLMs are *not* agentic. They do not have the judgement needed to decide if they have *achieved a goal* nor the judgement needed to assess if a task decomposition is appropriate for achieving a goal.

Assistant:

Thank you for sharing this outline - it's a crucial component of your argument that should definitely be addressed in the paper. Now I understand your position much more clearly: these systems aren't truly agentic because they fundamentally lack the judgment required for goal assessment and appropriate task decomposition.

Let me revise the introduction outline to incorporate this critique:

I've revised the introduction outline to properly incorporate your critique of agency claims. Now in the introduction, we'll:

1. Start by directly challenging OpenAI's "agentic capability" claims
2. Present Zvi's "mundane utility" concept more accurately as allowing focus on practical applications while setting aside larger questions
3. Explicitly outline why these systems lack true agency (can't assess goal achievement, can't evaluate task appropriateness, etc.)
4. Connect this lack of agency to the technoscholasticism concept

With these revisions, the introduction will more accurately set up your key arguments. Does this better reflect your position on agency and its relation to the rest of your critique?

User:

metadata: BB

Great, we'll want to footnote Zvi, to make sure that we note that the larger questions are existentially important -- they are merely outside the scope of the paper.

Also "verification" is a problem -- (I'm a philosopher of science, so that's a major loaded word. I think I want to nod in a more Kuhn/Lakatos direction rather than Carnap/Popper positivism. So maybe "A digital scholasticism that privileges textual authorities over critical assessment?" (I don't love this, let's workshop this. I wonder if we could work in Lakatos' research programme in)

The research questions How well can an LLM, using A Deep Research workflow assist in performing research?

Assistant:

You're absolutely right on both counts.

For Zvi's concept, a footnote is the perfect approach to acknowledge that existential questions are indeed important while maintaining our focus on practical applications.

Regarding "verification" - I see why that's problematic given your philosophy of science background. The positivist connotations wouldn't align with your perspective. Let me suggest some Kuhn/Lakatos-inspired alternatives:

- "A digital scholasticism that privileges textual authorities over critical assessment of knowledge claims"
- "A digital scholasticism that privileges textual authorities over confrontation with anomalies" (more Kuhnian)
- "A digital scholasticism that privileges textual authorities over engagement with competing research programmes" (more Lakatosian)

I particularly like the Lakatosian framing since it acknowledges that research involves evaluating competing explanatory frameworks rather than simple verification against facts.

For the research question, your simplified version is much clearer: "How well can an LLM, using a Deep Research workflow, assist in performing research?" We could also add a follow-up: "What are the epistemological limitations of such assistance?"

User:

metadata: BB

"A digital scholasticism that privileges textual authorities over critical assessment of knowledge claims"

As much as I love the Lakatosian one, this is more clear to most of the audience. And means I don't need to go into details about hard cores and protective belts in the *intro*.

2 / 2

Assistant:

I agree, that formulation is clearer and more accessible while still capturing the epistemological issue. Let me update the outline with this definition:

The outline has been updated with all your specifications:

1. Added a footnote indicator for Zvi's concept to acknowledge existential concerns
2. Updated the definition of technoscholasticism to "privileges textual authorities over critical assessment of knowledge claims"
3. Replaced the multiple research questions with your clearer single question about LLMs assisting in research

With these changes, the outline now better reflects your philosophical orientation and provides a clearer framework for drafting the introduction. Shall I proceed with writing the full introduction based on this outline?

User:

metadata: BB

Yes, let's see how it goes.

Assistant:

I've drafted the introduction based on our outline, incorporating all the key elements we discussed. The draft:

1. Opens with the OpenAI quote about Deep Research as a hook
2. Immediately challenges the "agentic capability" claims
3. Introduces Mowshowitz's "mundane utility" concept with the appropriate footnote
4. Presents your central thesis about judgment deficiency and introduces technoscholasticism
5. Explicitly defines technoscholasticism and connects it to historical scholasticism via the Horden citation
6. Addresses Dwarkesh's question through the technoscholasticism framework
7. Outlines your autoethnographic approach and the paper's contributions

The introduction is approximately 750 words as requested. I've maintained your academic register while ensuring clarity and theoretical depth. All citations use the BibLaTeX keys you provided.

Would you like me to make any adjustments to this draft?

User:

metadata: BB

Ok, after editing for a while, here's what I have as the first section.

(Note that I've made the language a touch more formal, Australian, and split a long paragraph)

Critically assess this as the paper's intro. Does it serve its purpose in the context of the rest of the plans? What other edits do we need to make before we can proceed?

Do we think the final parts of this are sufficient signposting, or should we add a % TODO SIGNPOST for later?

Assistant:

Your edited introduction is excellent - it shows notable improvements in several areas:

1. **Philosophical precision:** Adding "*qua* goal-following systems" and the more precise characterization of technoscholasticism grounds the argument firmly in philosophical tradition while introducing your novel theoretical contribution.
2. **Evidence of judgment failure:** The observation that LLMs never admit inadequacy provides concrete evidence of judgment deficiency, not just theoretical claims.
3. **Scholarly tone:** The Australian spelling and more formal phrasing establish appropriate academic register, while the "plastic research pal" reference adds a scholarly touch of humor that works well.
4. **Expanded footnote:** The additional context on existential questions shows nuanced understanding of the broader implications while maintaining focus.
5. **Carnapian reference:** The stronger positioning against Carnapian confirmation helps situate your argument in philosophy of science traditions.

Areas for consideration:

The final paragraph's signposting is thematically sound but could benefit from more specific structural guidance. I recommend adding a % TODO SIGNPOST marker to revisit this after other sections are developed. This will ensure the introduction properly maps to the paper's final structure.

Overall, this introduction successfully establishes:

- The problem (false agency claims)
- Your thesis (judgment deficiency explained through technoscholasticism)
- Your methodological approach (autoethnographic)
- The paper's contribution (realistic assessment of limitations and applications)

It serves its purpose well in setting up the technoscholasticism framework that will drive the rest of the paper.

User:

metadata: BB

Ok, this chat is getting a touch long. There's a chance that I'm going to have to refresh.

Thus, let's establish a context that we can pass to future-you, plus a checklist of documents that I should attach to the chat.

Please make an *exhaustively detailed* context that we can pass to the future instance of you. (Also, please indicate if I should fully copy-paste this chat. If not, please quote EXTENSIVELY) such that future-you can preserve most of the necessary context for the next section. (There's a chance we don't need to do this, but now's a good time for a stocktake.)

Remember, I'm going to be attaching the document. If we're talking about citations, we need to be talking about the citations we *have not yet used* in the text. Specifically, we want to talk about AND QUOTE (this should probably be an artifact) all of the ideas that we have for all of the upcoming sections, to better detail the outline (which we will also be supplying.)

2 / 2

Assistant:

</thinking>

I've created an exhaustively detailed context summary to pass to future versions of me. This artifact includes:

1. Current progress on the introduction
2. Key theoretical framework elements (technoscholasticism, judgment deficiency, agency critique)
3. Essential documents needed for future sessions
4. **Extensive quotes from transcripts** organized by section relevance
5. Unused citations that would be valuable for upcoming sections
6. Next steps with detailed outline for Section 2

I recommend copying the complete chat history if possible, as this will provide the richest context. If that's not feasible, this artifact contains the most essential information and quotes needed for continuity.

For our next session, we should focus on drafting Section 2 (Theoretical Framework) following the outline provided. We should continue our approach of first creating a detailed sub-paragraph outline, refining it based on your feedback, and then drafting the full section.

Would you like me to make any modifications or additions to this context summary before we conclude this session?

User:

metadata: BB

Great. And can you also generate a roughly 100 word abstract for this entire document? I want to present my work to an upcoming conference.

Assistant:

</thinking>

Here's an abstract for your paper that captures the core contribution and findings in approximately 100 words:

The abstract highlights your novel theoretical contribution ("technoscholasticism"), the methodological approach (autoethnographic investigation), the three dimensions of judgment deficiency, and both the theoretical and practical implications of your work. I've used Australian English spelling and maintained the formal academic register established in your introduction.

This should provide conference reviewers with a clear understanding of your paper's significance in critically examining AI research tools while proposing constructive alternatives that acknowledge their limitations.

User:

metadata: BB

Given all of these sections, give me a part of the outline to talk about into my phone recorder as I walk to work. Preferably lower in the outline. What needs more detail for you to work from?

Assistant:

For your walk to work, I'd recommend focusing on **Section 6: Research Tool Design Requirements**. This constructive section needs more specific details about how research tools should be designed to compensate for AI judgment deficiencies.

Specifically, please elaborate on subsection 6.2 (Research Process Support):

User:

\subsection{Research Process Support}

- Stage-based design (literature finding, extraction, etc.)

- Non-waterfall approach to research flows
- Structured to-do lists with human checkpoint integration
- Outline-driven context switching
- Documentation of decision points

Consider addressing these questions in your dictation:

1. What would an effective stage-based design look like in practice? How should transitions between stages work?
2. You've mentioned that research can't follow a waterfall SDLC - what alternative workflow model would be more appropriate? How would it accommodate iteration and refinement?
3. How exactly should human checkpoints be integrated? At what specific points in the research process is human judgment most critical?
4. How should outline-driven context switching work in practice? How would the system know which sources to load based on outline position?
5. What documentation of decision points would be most valuable for researchers reviewing an AI's process?

This section bridges your theoretical critique with practical recommendations, making it particularly valuable for the paper's contribution.

User:

metadata: BB

Great. I've recorded answers to all of those in another chat. Give me more questions on a different section to focus on next, please?

Assistant:

For your next recording session, I'd recommend focusing on **Section 5: Analysis Through Judgment Framework**, particularly the subsection on manifestations of technoscholasticism:

User:

- ```
\subsection{Manifestations of Technoscholasticism}
```
- Appeals to authority without verification
  - Absence of empirical checks
  - Conflation of description with insight
  - Inability to identify true knowledge gaps

Consider addressing these questions in your dictation:

1. What specific examples from your autoethnographic testing best demonstrate technoscholasticism in action? Can you describe instances where AI tools privileged textual authority over critical assessment?
2. How did you observe the AI systems conflating description with insight? What markers distinguish genuine insight from sophisticated description?
3. In what ways did the AI tools fail to identify true knowledge gaps in the literature? Did they simply aggregate existing views or did they attempt to highlight contradictions?
4. Did you observe any instances where the AI tools appeared to challenge their technoscholastic tendencies? Were there conditions under which they performed better?
5. How does the manifestation of technoscholasticism differ across different tool types (Deep Research vs. Computer Use vs. Elicit)?

This section is crucial for connecting your theoretical framework to empirical observations, providing evidence for your critique while acknowledging nuances in AI performance.

---

**User:**

metadata: BB

**OK. Questions for section 3**

---

**Assistant:**

For **Section 3: Methodology**, I recommend addressing these questions to strengthen the methodological foundation of your paper:

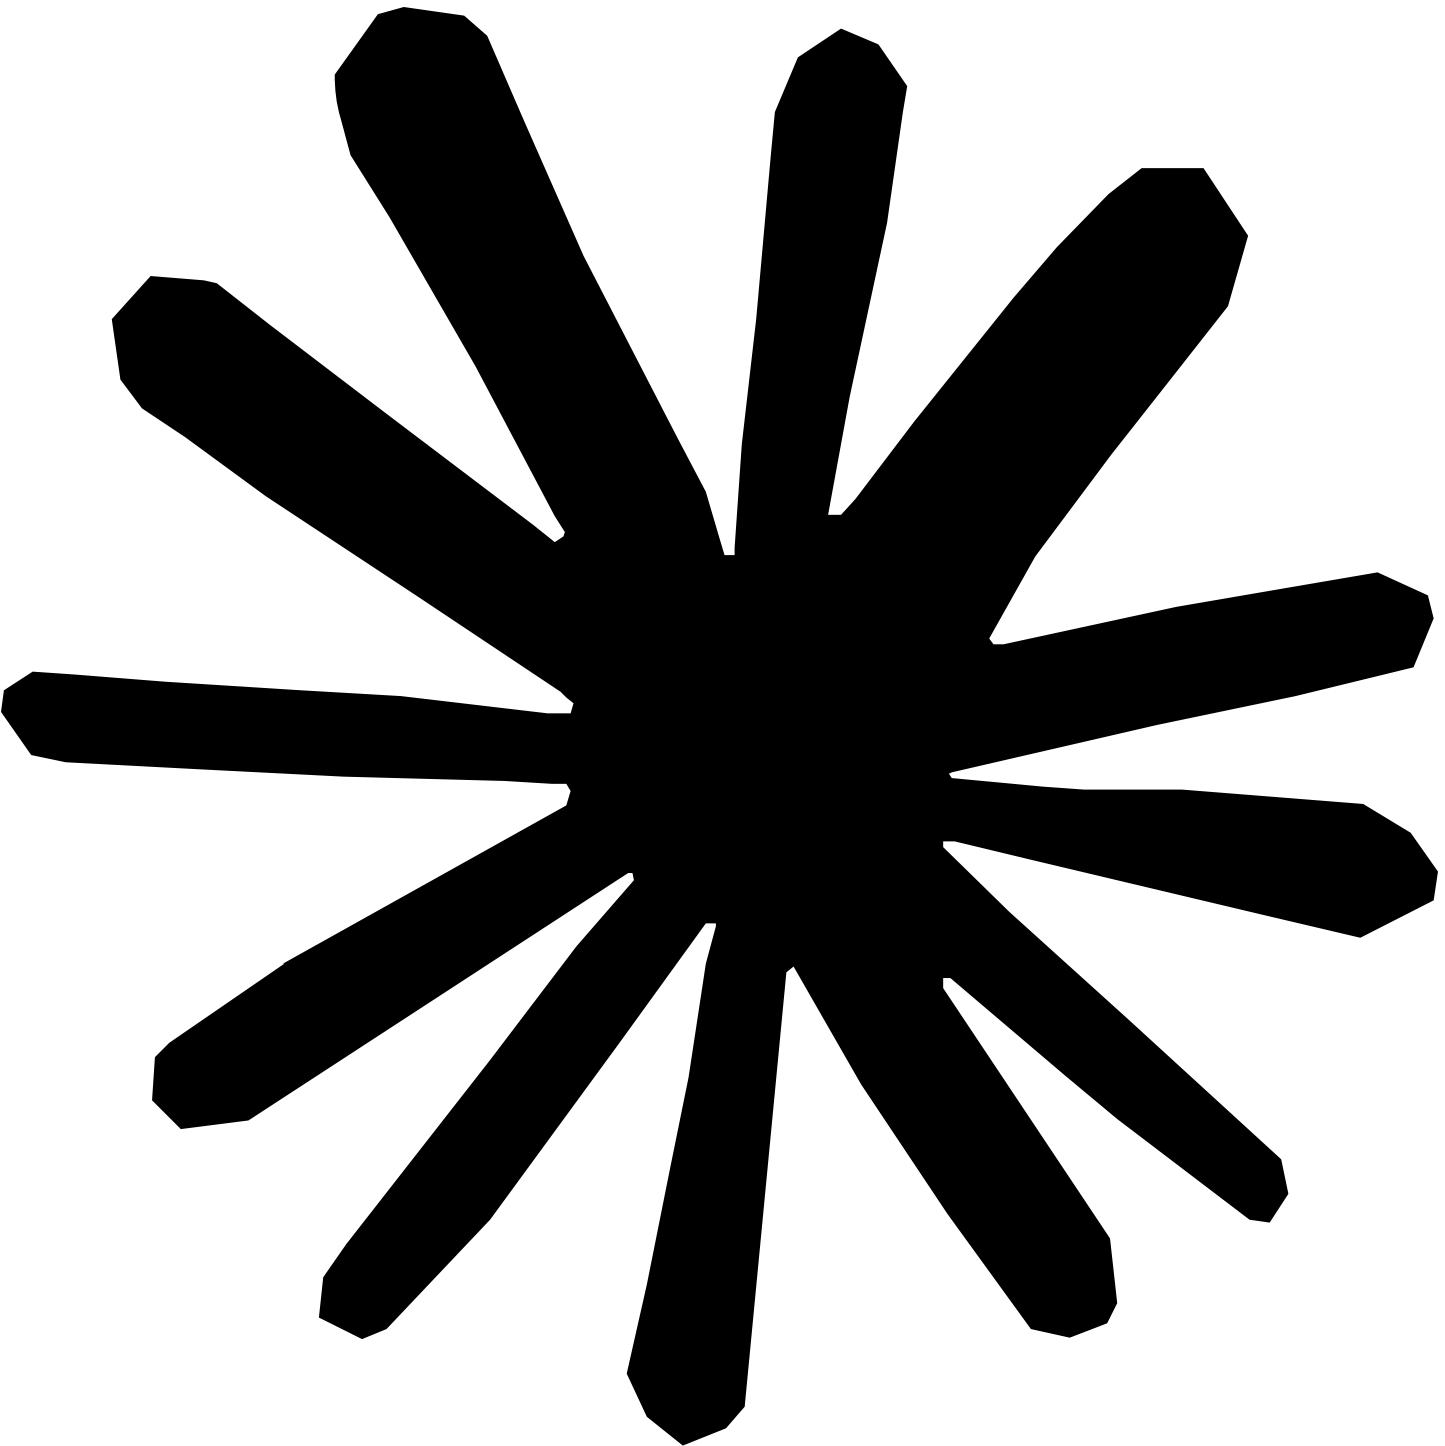
---

**User:**

```
\section{Methodology}
\subsection{Autoethnographic Investigation}
```

1. Why is autoethnography particularly appropriate for studying AI research tools? How does it capture dimensions that other methodological approaches might miss?
2. What specific autoethnographic practices did you employ during your testing sessions (e.g., reflective journaling, recorded dialogues, iteration documentation)? How were these materials analyzed?
3. How did you balance your role as both participant and observer during the testing process? What steps did you take to mitigate potential biases?
4. What epistemological position underpins your autoethnographic approach? How does this align with your critique of technoscholasticism?
5. How did you establish trustworthiness in your autoethnographic data? What criteria should readers use to evaluate the validity of your observations?
6. How does your approach differ from similar studies examining AI tools? What precedents exist for using autoethnography in studying AI capabilities?

These questions will help you articulate a robust methodological framework that anticipates potential critiques while establishing the scientific credibility of your autoethnographic approach.



[Claude can make mistakes.](#)  
[Please double-check responses.](#)

No file chosen

This conversation has reached its [maximum length](#). You may want to remove some files from your project's knowledge base.

Ok.

---

**metadata:**

**User:**

**Chat  
controls**

## User:

Claude  
Sonnet

4

Smart,  
efficient  
model  
for  
everyday  
use [Learn  
more](#)

## Artifacts

- 
- 
- 
- 
- 

## Project content

 [AbsenceJudgement](#)

Created  
by  
Brian  
Ballsun-  
Stanton

**Denubis/LLM-  
History-  
Paper**

main



**Content** [GITHUB](#)

**Denubis/LLM-  
History-  
Paper**

main



[GITHUB](#)